

## Assignment 7 for STA 601

(Richard) Fangjian Guo

NetID: fg47

Department of Computer Science

September 24, 2013

Let us denote the test set with  $(y_1, y_2)$  and the training set  $(y_1^o, y_2^o)$ .

### 1 Oracle prediction

The oracle predictor is simply

$$\hat{y}_1^{(\text{oracle})} = E[y_1 | y_2] = \mu_{1|2},$$

where

$$y_1 | y_2 \sim N(\mu_{1|2}, \sigma_{1|2}^2).$$

And the predictor is computed with true model parameters  $\mu$  and  $\Sigma$ .

### 2 MLE prediction

The MLE predictor is again

$$\hat{y}_1^{(\text{MLE})} = E(y_1 | y_2) = \hat{\mu}_{1|2},$$

where  $\mu_{1|2}$  is estimated with training set data  $(y_1^o, y_2^o)$ , namely

$$\hat{y}_1^{(\text{MLE})} = \hat{\mu}_{1|2} = \hat{\mu}_1 + \hat{\Sigma}_{12}(\hat{\Sigma}_{22}^{-1})(y_2 - \hat{\mu}_2).$$

$\hat{\mu}$  and  $\hat{\Sigma}$  are estimated from the training set as sample mean and sample covariance respectively.

### 3 Bayesian prediction

The Bayesian predictor is the mean for the posterior predictive distribution for each  $y_2$ , i.e.

$$\hat{y}_1^{(\text{Bayes})} = E(y_1 | y_2, \{y_1^o, y_2^o\}).$$

This predictor can be computed with the following procedure.

For each  $y_2$ , and for each iteration  $s = 1, 2, \dots, S$ ,

1. Sample  $\mu^{(s)}$  and  $\Sigma^{(s)}$  from the posterior distribution with Gibbs sampling.

2. Compute the parameters for conditional distribution with

$$\mu_{1|2}^{(s)} = \mu_1^{(s)} + \Sigma_{12}^{(s)} (\Sigma_{22}^{(s)})^{-1} (y_2 - \mu_2^{(s)}),$$

and

$$\sigma_{1|2}^{2(s)} = \Sigma_{11}^{(s)} - (\Sigma_{12}^{(s)})^2 (\Sigma_{22}^{(s)})^{-1}.$$

3. Sample  $y_1^{(s)}$  from the conditional distribution  $N(\mu_{1|2}^{(s)}, \sigma_{1|2}^{2(s)})$ .

Then the predictor is estimated as their average, namely

$$\hat{y}_1^{(\text{Bayes})} = \frac{1}{S} \sum_{s=1}^S y_1^{(s)}.$$

The predictive interval is also estimated from the quantiles of these samples.

#### 4 Results

The Bayesian predictor with 95% predictive interval, aside with the test set, is given by Figure 1.




Figure 1: The Bayesian predictor with 95% predictive interval, compared with the test set.

As can be noted in the figure, the coverage rate of the predictive intervals is 100%.

The MSE (Mean Square Error) for Bayesian, MLE and oracle predictors is summarized in Table 1. Bayesian is slightly better than MLE predictor. The three predictors are compared to the test set data in Figure 2.

#### 5 Comparing MLE and linear regression

The two models, despite their different definitions, actually result in the same predictor function, as shown in Figure 3. This is because

	Bayesian	MLE	Oracle
MSE	0.36949	0.37238	0.36908

Table 1: MSE

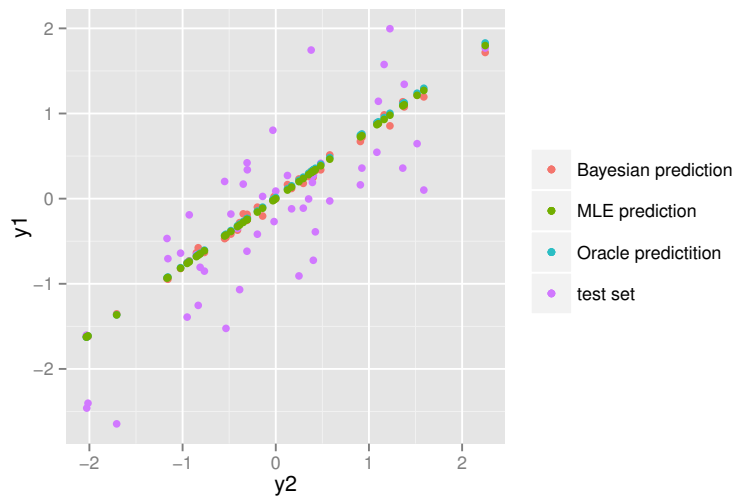


Figure 2: The predicted values compared with test set.

the close form solution to linear regression is the same as the MLE predictor.




Figure 3: The MLE predictor compared with linear regression line.

## 6 Appendix: R code

```
1 library(ggplot2)
```

```

2 library(reshape)
3 library(mvtnorm)
4 library(monomvn)
5
6 # training set
7 mu <- c(0,0)
8 sigma <- matrix(c(1,0.8,0.8,1), nrow=2, ncol=2, byrow=T)
9 y.sp <- rmvnorm(100, mu, sigma)
10
11 # new instances ( test set )
12 n <- 50
13 y.new <- rmvnorm(n, mu, sigma)
14
15 # oracle predictions
16 predict.oracle <- mu[1] + sigma[1,2]/sigma[2,2] * (y.new[,2] - mu[2])
17 MSE.oracle <- mean((predict.oracle - y.new[,1])^2)
18
19 # MLE estimates
20 mu.mle <- apply(y.sp, 2, mean)
21 sigma.mle <- cov(y.sp)
22 ## MLE gaussian conditional expectation
23 predict.mle <- mu.mle[1] + sigma.mle[1,2]/sigma.mle[2,2] * (y.new[,2] - mu.mle[2])
24 MSE.mle <- mean((predict.mle - y.new[,1])^2)
25
26 # Bayesian
27 # hyperpara
28 mu.o <- c(0,0)
29 Lambda.o <- matrix(c(1,0.5,0.5,1), nrow=2, byrow=T)
30 S.o <- Lambda.o
31 nu.o <- 4
32 # consts
33 y.mean <- mu.mle
34 ## Gibbs sampler
35 ## posterior predictive for each y.new[,2]
36 B <- 500
37 S <- 500
38 sigma.run <- sigma.mle
39 y1.predict.bayes.mat <- NULL
40 for (t in 1:(B+S)) {
41   Lambda.run <- solve(solve(Lambda.o) + n*solve(sigma.run))
42   mu.run <- Lambda.run %*% (n*solve(sigma.run) %*% y.mean) # mu.o is zero here
43   theta.run <- rmvnorm(1, mu.run, Lambda.run)
44   S.run <- S.o + (t(y.sp) - c(theta.run)) %*% t((t(y.sp) - c(theta.run)))
45   sigma.run <- solve(rwish(nu.o+n, solve(S.run)))
46   # burn-in
47   if (t <= B) {
48     next
49   }
50   # posterior predictive
51   # parameters for each new y2
52   sigma1.conditional <- (sigma.run[1,1] - sigma.run[1,2]/sigma.run[2,2]*sigma.run
53     [1,2]) * rep(1,n)
54   mu1.conditional <- theta.run[1] + sigma.run[1,2]/sigma.run[2,2]*(y.new[,2] - theta.
55     run[2])
56   y1.tmp <- numeric(n)
57   # sampling conditional
58   for (j in 1:n) {
59     y1.tmp[j] <- rnorm(1, mu1.conditional[j], sqrt(sigma1.conditional[j]))
60   }
61   y1.predict.bayes.mat <- rbind(y1.predict.bayes.mat, y1.tmp)
62 }

```

```
61 # prediction
62 predict.bayes.mean <- apply(y1.predict.bayes.mat, 2, mean)
63 predict.bayes.interval . left <- apply(y1.predict.bayes.mat, 2, function(x) quantile(x,
64 probs=c(0.025)))
65 predict.bayes.interval . right <- apply(y1.predict.bayes.mat, 2, function(x) quantile(x,
66 probs=c(0.975)))
67 MSE.bayes <- mean((predict.bayes.mean - y.new[,1])^2)
68 interval . hit <- sum(as.numeric(predict.bayes.interval.left <= y.new[,1] & y.new[,1]
69 <= predict.bayes.interval.right))/n
70
71 # plotting
72 ## plot
73 predict.df <- data.frame(y1=y.new[,1], y2=y.new[,2],
74 predict.mle, predict.oracle, predict.bayes.mean,
75 predict.bayes.interval . left, predict.bayes.interval . right)
76 fig1 <- ggplot(predict.df) +
77   geom_point(aes(x=y2, y=predict.bayes.mean, color="Bayesian_prediction")) +
78   geom_point(aes(x=y2, y=y1, color="test_set")) +
79   geom_point(aes(x=y2, y=predict.mle, color="Oracle_prediction")) +
80   geom_point(aes(x=y2, y=predict.oracle, color="MLE_prediction")) +
81   theme(legend.title=element_blank()) + ylab("y1")
82 print(fig1)
```

---